



Keras Tuner hyperparameter search (996 trials)


Hyperparameter sensitivity




Hidden layers rho = +0.45




Trainable params rho = +0.41



L2 reg. rho = +0.20



Dropout rate rho = +0.19



● Trial ★ Best trial (0.0010) — Binned median